

RAINFALL

June 15, 2023

#PROBLEM STATEMENT:- TO PREDICT THE RAIN FALL BASED ON VARIOUS FEATURES OF THE DATASET WHICH IS GIVEN BELOW

IMPORTING THE ESSENTIAL LIBRARIES:-

```
[3]: import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[5]: df=pd.read_csv(r"/content/RAIN FALL.csv")
df
```

```
[5]:
```

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	\	
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5		
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1		
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9		
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1		
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7		
...		
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6		
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0		
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2		
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1		
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6		
	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	\
0	365.1	481.1	332.6	388.5	558.2	33.6	3373.2	136.3	560.3	
1	228.9	753.7	666.2	197.2	359.0	160.5	3520.7	159.8	458.3	
2	728.4	326.7	339.0	181.2	284.4	225.0	2957.4	156.7	236.1	
3	502.0	160.1	820.4	222.2	308.7	40.1	3079.6	24.1	506.9	
4	368.7	330.5	297.0	260.7	25.4	344.7	2566.7	1.3	309.7	
...	
4111	350.2	254.0	255.2	117.4	184.3	14.9	1533.7	7.9	196.2	
4112	231.5	381.2	179.8	145.9	12.4	8.8	1405.5	19.3	99.6	
4113	296.4	154.4	180.0	72.8	78.1	26.7	1426.3	60.6	131.1	

4114	116.1	466.1	132.2	169.2	59.0	62.3	1395.0	69.3	76.7
4115	257.5	146.4	160.4	165.4	231.0	159.0	1642.9	2.7	223.9

	Jun-Sep	Oct-Dec
0	1696.3	980.3
1	2185.9	716.7
2	1874.0	690.6
3	1977.6	571.0
4	1624.9	630.8
...
4111	1013.0	316.6
4112	1119.5	167.1
4113	1057.0	177.6
4114	958.5	290.5
4115	860.9	555.4

[4116 rows x 19 columns]

DATA PREPROCESSING:-

```
[6]: df.head()
```

```
[6]:
```

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	\
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	

	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	\
0	365.1	481.1	332.6	388.5	558.2	33.6	3373.2	136.3	560.3	
1	228.9	753.7	666.2	197.2	359.0	160.5	3520.7	159.8	458.3	
2	728.4	326.7	339.0	181.2	284.4	225.0	2957.4	156.7	236.1	
3	502.0	160.1	820.4	222.2	308.7	40.1	3079.6	24.1	506.9	
4	368.7	330.5	297.0	260.7	25.4	344.7	2566.7	1.3	309.7	

	Jun-Sep	Oct-Dec
0	1696.3	980.3
1	2185.9	716.7
2	1874.0	690.6
3	1977.6	571.0
4	1624.9	630.8

```
[7]: df.tail()
```

```
[7]:
```

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	\
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	

4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4

	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	Jun-Sep	Oct-Dec
4111	255.2	117.4	184.3	14.9	1533.7	7.9	196.2	1013.0	316.6
4112	179.8	145.9	12.4	8.8	1405.5	19.3	99.6	1119.5	167.1
4113	180.0	72.8	78.1	26.7	1426.3	60.6	131.1	1057.0	177.6
4114	132.2	169.2	59.0	62.3	1395.0	69.3	76.7	958.5	290.5
4115	160.4	165.4	231.0	159.0	1642.9	2.7	223.9	860.9	555.4

```
[8]: df.isnull().any()
```

```
[8]: SUBDIVISION    False
YEAR              False
JAN               True
FEB               True
MAR               True
APR               True
MAY               True
JUN               True
JUL               True
AUG               True
SEP               True
OCT               True
NOV               True
DEC               True
ANNUAL            True
Jan-Feb           True
Mar-May           True
Jun-Sep           True
Oct-Dec           True
dtype: bool
```

```
[9]: df.fillna(method='ffill',inplace=True)
```

```
[10]: df.describe()
```

```
[10]:
```

	YEAR	JAN	FEB	MAR	APR \
count	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000
mean	1958.218659	18.957240	21.823251	27.415379	43.160641
std	33.140898	33.576192	35.922602	47.045473	67.816588
min	1901.000000	0.000000	0.000000	0.000000	0.000000
25%	1930.000000	0.600000	0.600000	1.000000	3.000000
50%	1958.000000	6.000000	6.700000	7.900000	15.700000
75%	1987.000000	22.200000	26.800000	31.400000	50.125000

max	2015.000000	583.700000	403.500000	605.600000	595.100000
-----	-------------	------------	------------	------------	------------

	MAY	JUN	JUL	AUG	SEP \
count	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000
mean	85.788994	230.567979	347.177235	290.239796	197.524781
std	123.220150	234.896056	269.321089	188.785639	135.509037
min	0.000000	0.400000	0.000000	0.000000	0.100000
25%	8.600000	70.475000	175.900000	155.850000	100.575000
50%	36.700000	138.900000	284.800000	259.400000	174.000000
75%	97.400000	306.150000	418.325000	377.800000	266.225000
max	1168.600000	1609.900000	2362.800000	1664.600000	1222.000000

	OCT	NOV	DEC	ANNUAL	Jan-Feb \
count	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000
mean	95.724198	40.081997	19.042225	1417.221769	40.768975
std	99.689878	68.851397	42.655830	907.547328	59.302112
min	0.000000	0.000000	0.000000	62.300000	0.000000
25%	14.600000	0.700000	0.100000	806.450000	4.100000
50%	65.750000	9.700000	3.100000	1124.650000	19.200000
75%	148.600000	46.325000	17.600000	1660.425000	50.425000
max	948.300000	648.900000	617.500000	6331.100000	699.500000

	Mar-May	Jun-Sep	Oct-Dec
count	4116.000000	4116.000000	4116.000000
mean	156.579155	1065.552114	154.957070
std	202.056770	707.840186	167.807169
min	0.000000	57.400000	0.000000
25%	24.200000	574.375000	34.200000
50%	75.150000	881.750000	98.800000
75%	197.700000	1291.125000	215.775000
max	1745.800000	4536.900000	1252.500000

```
[11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4116 entries, 0 to 4115
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype
---  -
0   SUBDIVISION     4116 non-null   object
1   YEAR            4116 non-null   int64
2   JAN             4116 non-null   float64
3   FEB             4116 non-null   float64
4   MAR             4116 non-null   float64
5   APR             4116 non-null   float64
6   MAY             4116 non-null   float64
7   JUN             4116 non-null   float64
```

```

8    JUL          4116 non-null    float64
9    AUG          4116 non-null    float64
10   SEP          4116 non-null    float64
11   OCT          4116 non-null    float64
12   NOV          4116 non-null    float64
13   DEC          4116 non-null    float64
14   ANNUAL       4116 non-null    float64
15   Jan-Feb      4116 non-null    float64
16   Mar-May      4116 non-null    float64
17   Jun-Sep      4116 non-null    float64
18   Oct-Dec      4116 non-null    float64
dtypes: float64(17), int64(1), object(1)
memory usage: 611.1+ KB

```

```
[12]: df.columns
```

```
[12]: Index(['SUBDIVISION', 'YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL',
          'AUG', 'SEP', 'OCT', 'NOV', 'DEC', 'ANNUAL', 'Jan-Feb', 'Mar-May',
          'Jun-Sep', 'Oct-Dec'],
          dtype='object')
```

```
[13]: df.shape
```

```
[13]: (4116, 19)
```

```
[14]: df['ANNUAL'].value_counts()
```

```
[14]: 790.5      4
      770.3      4
      1836.2     4
      1024.6     4
      1926.5     3
      ..
      443.9      1
      689.0      1
      605.2      1
      509.7      1
      1642.9     1
      Name: ANNUAL, Length: 3712, dtype: int64
```

```
[15]: df['Jan-Feb'].value_counts()
```

```
[15]: 0.0      238
      0.1      80
      0.2      52
      0.3      38
      0.4      32
```

```

...
23.3      1
95.2      1
76.9      1
66.5      1
69.3      1
Name: Jan-Feb, Length: 1220, dtype: int64

```

```
[16]: df['Mar-May'].value_counts()
```

```

[16]: 0.0      29
      0.1      13
      0.3      11
      8.3      11
     11.5      10
      ..
    246.3      1
    248.1      1
    151.3      1
    249.5      1
    223.9      1
Name: Mar-May, Length: 2262, dtype: int64

```

```
[17]: df['Jun-Sep'].value_counts()
```

```

[17]: 434.3      4
      334.8      4
      573.8      4
      613.3      4
     1082.3      3
      ..
    301.6      1
    380.9      1
    409.3      1
    229.4      1
    958.5      1
Name: Jun-Sep, Length: 3683, dtype: int64

```

```
[18]: df['Oct-Dec'].value_counts()
```

```

[18]: 0.0      16
      0.1      15
      0.5      13
      0.6      12
      0.7      11
      ..
    191.5      1

```

```

124.5    1
139.1    1
41.5     1
555.4    1
Name: Oct-Dec, Length: 2389, dtype: int64


```

EXPLORATORY DATA ANALYSIS:-

```

[19]: df=df[['JAN','FEB','MAR','APR','DEC']]
      sns.heatmap(df.corr(),annot=True)
      plt.show()

```



```

[20]: df.columns

```

```

[20]: Index(['JAN', 'FEB', 'MAR', 'APR', 'DEC'], dtype='object')

```

```

[21]: x=df[["FEB"]]
      y=df[["JAN"]]

```

LINEAR REGRESSION:-

```
[22]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.
↪33,random_state=42)
```

```
[23]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(X_train,y_train)
print(reg.intercept_)
coeff_=pd.DataFrame(reg.coef_,x.columns,columns=['coefficient'])
coeff_
```

9.650666612303553

```
[23]:      coefficient
FEB      0.442278
```


```
[24]: score=reg.score(X_test,y_test)
print(score)
```

0.1793580786264921


```
[25]: predictions=reg.predict(X_test)
```

```
[26]: plt.scatter(y_test,predictions)
```

```
[26]: <matplotlib.collections.PathCollection at 0x7efc0260e1d0>
```


```
[27]: df500=df[:][:500]
sns.lmplot(x="FEB",y="JAN",order=2,ci=None,data=df500)
plt.show()
```



```
[28]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33)
      reg.fit(X_train,y_train)
      reg.fit(X_test,y_test)
```

```
[28]: LinearRegression()
```

```
[29]: y_pred=reg.predict(X_test)
      plt.scatter(X_test,y_test,color='black')
      plt.plot(X_test,y_pred,color='red')
      plt.show()
```



```
[30]: from sklearn.linear_model import LinearRegression
      from sklearn.metrics import r2_score
      model=LinearRegression()
      model.fit(X_train,y_train)
      y_pred=model.predict(X_test)
      r2=r2_score(y_test,y_pred)
      print("R2 Score:",r2)
```

R2 Score: 0.18612218989321283

RIDGE MODEL:-

```
[31]: from sklearn.linear_model import Lasso,Ridge
      from sklearn.preprocessing import StandardScaler
```

```
[75]: features= df.columns[0:5]
      target= df.columns[-5]
```

```
[76]: x=np.array(df['JAN']).reshape(-1,1)
      y=np.array(df['FEB']).reshape(-1,2)
```

```
[77]: x= df[features].values
      y= df[target].values
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.
↪3,random_state=17)
```

```
[78]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge=ridgeReg.score(x_train,y_train)
test_score_ridge=ridgeReg.score(x_test,y_test)
```


```
[79]: print("\n Ridge Model:\n")
print("the train score for ridge model is{}".format(train_score_ridge))
print("the test score for ridge model is{}".format(test_score_ridge))
```

Ridge Model:

```
the train score for ridge model is0.9999999999874192
the test score for ridge model is0.99999999998833
```

```
[80]: lr=LinearRegression()
```

```
[81]: plt.figure(figsize= (10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.
↪7,linestyle='none',marker="*",markersize=5,color='red',label=r'Ridge,$\alpha=10$',zorder=7)
plt.plot(features,alpha=0.
↪4,linestyle='none',marker='o',markersize=7,color="green",label='LinearRegression')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



LASSO MODEL:-

```
[82]: print("\n Lasso Model:\n")
lasso=Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls=lasso.score(x_train,y_train)
test_score_ls=lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is{}".format(test_score_ls))
```


Lasso Model:

The train score for ls model is 0.9999207747038827

The test score for ls model is 0.9999206791315255

```
[83]: pd.Series(lasso.coef_,features).sort_values(ascending=True).plot(kind="bar")
```

```
[83]: <Axes: >
```




```
[84]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,1,10],random_state=0).
        fit(x_train,y_train)
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

0.9999999999999921

0.9999999999999921

```
[85]: plt.figure(figsize= (10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.
        7,linestyle='none',marker="*",markersize=5,color='red',label=r'Ridge,$\alpha=10$',zorder=7)
plt.plot(lasso_cv.coef_,alpha=0.
        5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso,$\alpha=10$')
```

```
plt.plot(features,alpha=0.
↪4,linestyle='none',marker='o',markersize=7,color="green",label='LinearRegression')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



ELASTIC NET:-

```
[86]: from sklearn.linear_model import ElasticNet
elnet=ElasticNet()
elnet.fit(x,y)
print(elnet.coef_)
print(elnet.intercept_)
```

```
print(eln.score(x,y))
```

```
[9.99098574e-01 0.00000000e+00 3.02728910e-05 0.00000000e+00  
0.00000000e+00]  
0.016258606966612632  
0.9999992160905338
```

```
[88]: y_pred_elastic =eln.predict(x_train)  
mean_squared_error=np.mean((y_pred_elastic - y_train)**2)  
print(mean_squared_error)
```

```
0.0008816302333951295
```

#CONCLUSION:-

THE SCORE OF LINEAR REGRESSION IS :- 0.1793580786264921

THE SCORE OF RIDGE MODEL IS :- 0.99999999998833

THE SCORE OF LASSO MODEL IS :- 0.999999999999992

THE SCORE OF ELASTIC NET IS :- 0.9999992160905338

**-> AMONG ALL MODELS LASSO YEILD HIGHEST ACCURACY. SO, WE PRE-
FER LASSO MODEL FO R THIS DATA SET**